Highlights

- CPU and browser benchmark of three lightweight flood-segmentation UNets.
- On saturated FSSD, pretrained students match the teacher without distillation.
- INT8 quantization is architecture-dependent: Lite0 holds, two backbones collapse.
- FSSD models false-alarm on 70–100% of dry scenes; retraining cuts this to 3–9%.
- On harder data, distillation lifts flood IoU but copies teacher false alarms.